

LAB - 3

(24/3/25)

Simple Linear Regression

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def estimate_coef(x, y):

n = np.size(x)

m_x = np.mean(x)

m_y = np.mean(y)

ss_xy = np.sum((x - m_x) * (y - m_y))

ss_xx = np.sum((x - m_x) ** 2)

b_1 = ss_xy / ss_xx

b_0 = m_y - b_1 * m_x

return [b_0, b_1]

def plot(x, y, b):

plt.scatter(x, y, color='m')

y_pred = b[0] + b[1] * x

plt.plot(x, y_pred, color='g')

plt.xlabel('x')

plt.ylabel('y')

plt.show()

file_path = input("Enter path:")

df = pd.read_csv(file_path)

x = df.iloc[:, 0].values

y = df.iloc[:, 1].values

b = estimate_coef(x, y)

plot(x, y, b)

Output:

Enter path : /content/tvmarketing.csv

Estimated coefficients:

b_0 = 7.032

b_1 = 0.047

Multiple Linear Regression

```
data = {  
    "Feature1": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
    "Feature2": [2, 3, 5, 7, 11, 13, 17, 19, 23, 29],  
    "Feature3": [5, 6, 9, 12, 15, 18, 21, 24, 27, 30],  
    "Target": [5, 9, 15, 22, 31, 41, 53, 66, 80, 96]  
}
```

```
df = pd.DataFrame(data)
```

```
X = df.drop(columns=["Target"]).values
```

```
y = df["Target"].values.reshape(-1, 1)
```

```
X = np.hstack((np.ones((X.shape[0], 1)), X))
```

```
beta = np.linalg.solve(X.T @ X + 0.01 * np.identity(X.shape[1]), X.T @ y)
```

```
y_pred = X @ beta
```

```
mse = np.mean((y - y_pred) ** 2)
```

```
tot_var = np.sum((y - np.mean(y)) ** 2)
```

```
exp_var = np.sum((y_pred - np.mean(y)) ** 2)
```

```
r2 = exp_var / tot_var
```

```
print("Output:")
```

Output:

Model Coefficients: [0.040, 3.313, 0.120]

Intercept: -3.1599

Mean Squared Error: 5.362

R-squared score: 0.993

Logistic Regression

```
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))
```

```
def compute(X, y, theta):
```

```
    m = len(y)
```

```
    h = sigmoid(X @ theta)
```

```
    cost = (-1/m) * np.sum(y * np.log(h) + (1-y) * np.log(1-h))
```

```
    return cost
```

```
def grad-descent(X, y, theta, alpha, iterations):
```

```
    m = len(y)
```

```
    cost-history = []
```

```
    for _ in range(iterations):
```

```
        gradient = (1/m) * X.T @ (sigmoid(X @ theta) - y)
```

```
        theta -= alpha * gradient
```

```
    return theta, cost-history
```

```
def predict(X, theta):
```

```
    return (sigmoid(X @ theta) >= 0.5).astype(int)
```

```
accuracy = np.mean(y == pred)
```

```
print('accuracy: ', accuracy, 2/3)
```

Output:

Accuracy: 1